

Learning Docker



Table of contents


1. [What is Docker?](#)
2. [Why use Docker?](#)
3. [Docker Image](#)
4. [Dockerfile](#)
5. [Docker Compose](#)
6. [Volumes and Bind Mounts](#)
7. [Exposing the GPU](#)



What is Docker?

"Docker is an open platform for developing, shipping, and running applications."

"Docker provides the ability to package and run an application in a loosely isolated environment called a **container**. "



Why use Docker and Containers?

Self-contained: Has everything needed to run the application without relying on what is currently installed in the host:

- Allows us to have standardized environments.
- Easier to share and set up in new machines.
- Provides isolation and security.



Installing Docker

The github repo has three bash scripts:


- 0-Utills/install-docker
- 0-Utills/install-docker-compose
- 0-Utills/docker-post-installation

In the terminal, do:

```
chmod +x <script>
```

```
<script>
```

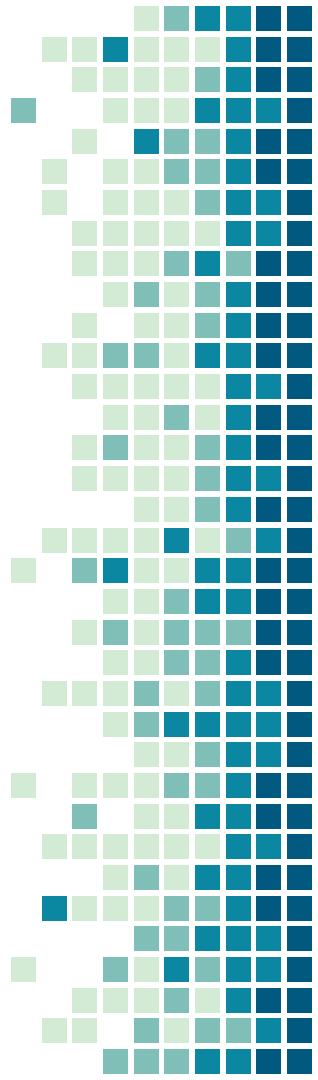
In VS Code: Add the Docker Extension



Docker Image

"An image is a **read-only template** with instructions for creating a Docker container. Often, an image is based on another image, with some additional customization."

Ex: an image based on ubuntu with conda installed.



Docker Image

Has instruction and files + metadata to create a system

- Files form the root filesystem
- Commands to execute when starting the container
- Environment variables

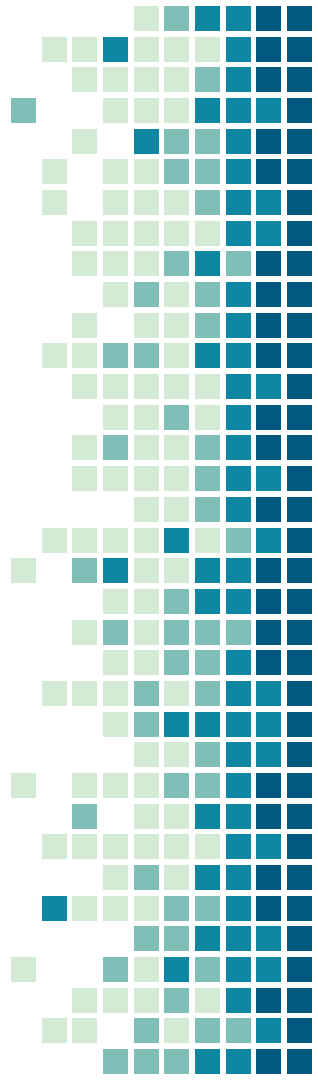


Docker Image

Multiple containers can be created based on an image.

An image is conceptually similar to a *Class*

A container would be similar to an *instance*

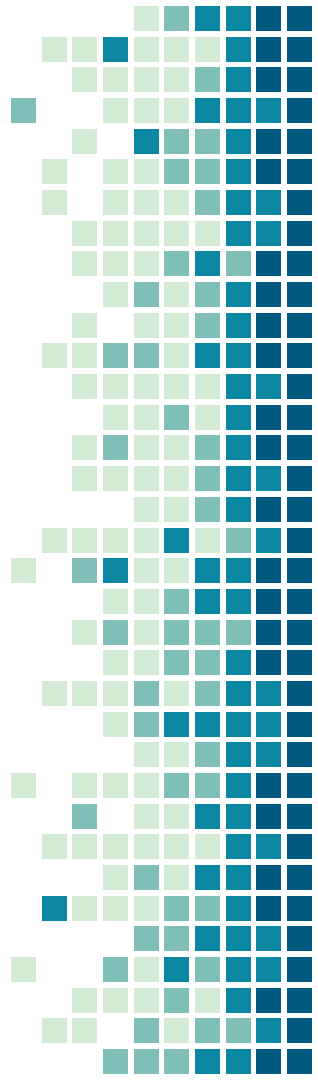


Docker Image

If an image is read-only, how do we change it?

We don't, we create a new image from the image of interest with all the changes we want.

How?

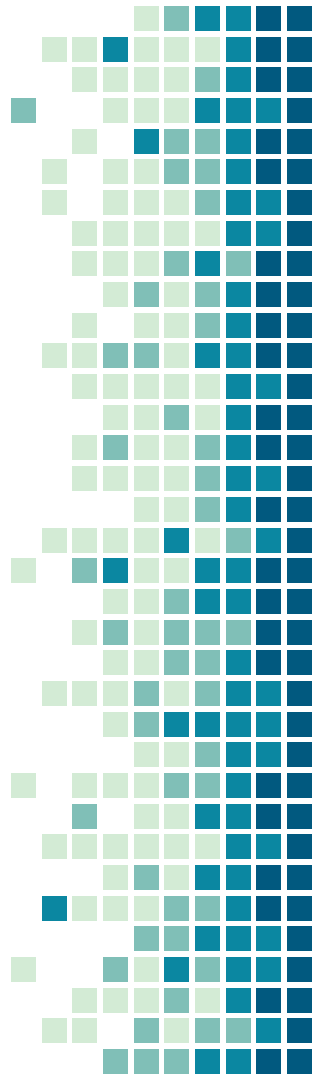


Dockerfile

Instructions to create an image. Each instruction in a Dockerfile creates a layer in the image.

Building images is a **fast and lightweight** virtualization process:

- When you change the Dockerfile and rebuild the image, only those layers which have changed are rebuilt



Dockerfile – Instructions

- **ARG** defines a variable that users can pass at build-time.
- **FROM** indicates the base image for our build.
- **WORKDIR** sets the working directory.
- **RUN** lines will be executed by Docker during the build.
 - Our RUN commands must be non-interactive



Dockerfile - Instructions

- **EXPOSE** informs Docker that the container listens on the specified network ports at runtime.
- **ENV** sets the environment variable <key> to the value <value>. These will persist when a container is run from the resulting image.

□ `ENV MY_NAME="Anna"`

Dockerfile - Instructions

- **COPY** copies new files or directories from `<src>` and adds them to the filesystem of the image at the path `<dest>`.
 - `COPY home /mydir/`
 - `src`: paths are interpreted as relative to the source of the context of the build.
 - `dest`: absolute path, or a path relative to `WORKDIR`
- If you need multiple files from your context, **COPY** them individually, rather than all at once. This ensures that each step's build cache is only invalidated (forcing the step to be re-run) if the specifically required files change.


Dockerfile – Instructions

- **ENTRYPOINT** sets the image's main command that will be executed everytime we run the container.

RUN vs **ENTRYPOINT**:

RUN is only executed once: When we are building the image and we can have multiple **RUN** commands.

ENTRYPOINT is executed everytime we run the container and it can appear only once in the Dockerfile.



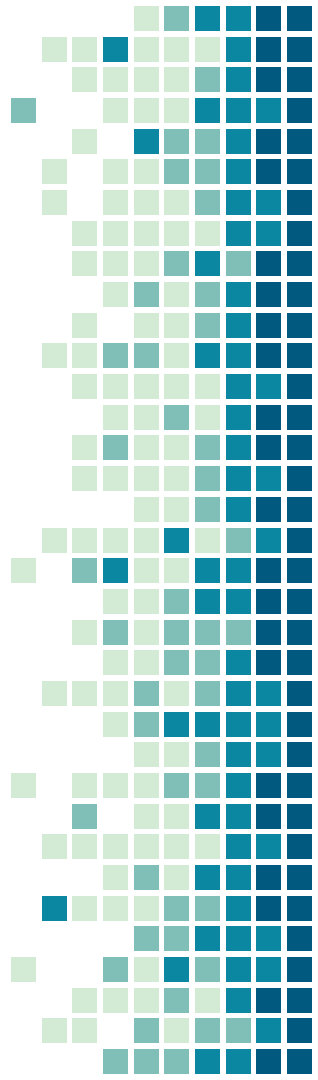
Dockerfile - Instructions

- **VOLUME** creates a mount point with the specified name and marks it as holding externally mounted volumes from native host.
 - We are going to learn more about volumes later.

Let's take a look at the example. (folder 1-Dockerfile)

Docker Compose

- Useful when:
 - We have multiple containers we want to connect or run together (e.g: Train, Re-train, Test, Predict). We will call them *services*.
 - We require particular run parameters.
- Not useful to:
 - Be used as an orchestrator.
 - Be used with Kubernetes.



Docker Compose – Services

Docker-compose files specify how each service will be configured. Some things we can configure are:

- **build** specifies the build configuration for creating the container image from source.
 - Dockerfile
 - Context
- **image** specifies the image to start the container from.



Docker Compose – Services

- `container_name` specifies a custom container name
- `volumes` defines mount host paths or named volumes that will be accessible by service containers. (more on [mounts and volumes](#) later)
- `command` overrides the the default command declared by the container image



Docker Compose – Services

- `deploy` specifies configuration for the deployment and lifecycle of services and runtime requirements. Some elements we can configure include:
 - `resources`: physical resources e.g. gpu
 - `cpus`: limit or reservation of cpu cores
 - `replicas`: number of containers that should be running at any given time.



Docker Volumes and Bind Mounts

Writable data for a container is ephemeral and it will be lost when the container is removed.

We can add persistent data to a container using *Docker volumes* or *Docker bind mounts*. They will store files on the local host filesystem.



Docker Volumes

- Managed by Docker
- Located in Docker's "private storage area."
(var/lib/docker/volumes/<volume ID>)
- Don't need to check directory exists before mounting it. Better for portability.
- No problems with permissions. Docker owns the volume.

Can be mounted to multiple containers

Docker Bind Mounts

- Managed by user/developer
- Directory can be located anywhere in the local system.
- Directory needs to exist on the local system. User will need to create it if it doesn't.
- Might encounter permission issues.

Docker Compose

Let's take a look at the example. (folder 2-Docker-compose)

Exposing the GPU

Make sure your local system has Nvidia drivers.


The github repo has two bash script for this:

- 0-Utills/install-nvidia-drivers
- 0-Utills/install-cuDNN (optional - read comments first)

In the terminal, do:

```
chmod +x <package>
```

```
<package>
```



Exposing the GPU

We need to configure our local system by installing the nvidia-container-runtime


The github repo has a bash script for this:

- 0-Utills/GPU-enable-for-Docker

In the terminal, do:

```
chmod +x 0-Utills/GPU-enable-for-Docker
```

```
0-Utills/GPU-enable-for-Docker
```



Exposing the GPU

Add device to Docker-compose

```
deploy:
```

```
  resources:
```

```
    reservations:
```

```
      devices:
```

```
        - driver: nvidia
```

```
          count: 1
```

```
          capabilities: [gpu]
```


Change Docker file to use image with GPU drivers/support

Exposing the GPU

Let's take a look at the example. (folder 3-Docker-GPU)



Questions?



Thank you

